

Part 1

Database on GCP Image

```
MySQL Shell 8.0.41

Copyright (c) 2016, 2025, Oracle and/or its affiliates.
Oracle is a registered trademark of Oracle Corporation and/or its affiliates.
Other names may be trademarks of their respective owners.

Type '\help' or '\?' for help; '\quit' to exit.
MySQL_JS> \sql
Switching to SQL mode... Commands end with ;
MySQL_SQL> \connect -h 104.197.129.116 -P 3306 -u jkim826
Creating a session to 'jkim826@104.197.129.116:3306'
Please provide the password for 'jkim826@104.197.129.116:3306': *****
Save password for 'jkim826@104.197.129.116:3306'? [Y]es/[N]o/[Ne]ver (default No): N
Fetching global names for auto-completion... Press ^C to stop.
Error during auto-completion cache update: Access denied; you need (at least one of) the PROCESS privilege(s) for this operation
Your MySQL connection id is 67521
Server version: 8.0.37-google (Google)
No default schema selected; type \use <schema> to set one.
MySQL_104.197.129.116:3306 ssl SQL> use music
Default schema set to 'music'.
Fetching global names, object names from 'music' for auto-completion... Press ^C to stop.
Error during auto-completion cache update: Access denied; you need (at least one of) the PROCESS privilege(s) for this operation
MySQL_104.197.129.116:3306 ssl music SQL> show tables
-> ;
+-----+
| Tables_in_music |
+-----+
| album_artists   |
| albums          |
| artists         |
| track_artists   |
| tracks          |
+-----+
5 rows in set (0.0314 sec)
MySQL_104.197.129.116:3306 ssl music SQL>
```

DDL Commands to Create Tables

- CREATE TABLE Album(AlbumId:INT [PK], AlbumPopularity:INT, AlbumName:VARCHAR(1000), TotalTracks:INT, AlbumDuration:INT)
- CREATE TABLE Artists(ArtistId:INT [PK], ArtistPopularity:INT, ArtistName:VARCHAR(1000))
- CREATE TABLE Songs (SongId:INT [PK], AlbumId:INT [FK], ArtistId:INT [FK], SongId:VARCHAR(1000), Explicit:INT, Energy:FLOAT, Danceability:FLOAT, Duration:INT)
- CREATE TABLE Streams(SongId:INT [PK][FK], Ranking:INT, NumStreams:INT)
- CREATE TABLE User(Username:VARCHAR(1000) [PK], Password:VARCHAR(1000), Email:VARCHAR(1000))
- CREATE TABLE UserTables(TableName:VARCHAR(1000) [PK], Username:VARCHAR(1000) [FK])

Images Showing At Least 1000 Rows in Three Different Tables Each

Albums Table:

```
MySQL 104.197.129.116:3306 ssl music SQL > SELECT COUNT(*)
-> FROM albums;

+-----+
| COUNT(*) |
+-----+
|      7470 |
+-----+
1 row in set (0.1426 sec)
```

Artists Table:

```
MySQL 104.197.129.116:3306 ssl music SQL > SELECT COUNT(*) FROM artists;

+-----+
| COUNT(*) |
+-----+
|      3812 |
+-----+
1 row in set (0.0533 sec)
```

Tracks Table:

```
MySQL 104.197.129.116:3306 ssl music SQL > SELECT COUNT(*) FROM tracks;

+-----+
| COUNT(*) |
+-----+
|      9993 |
+-----+
1 row in set (0.0444 sec)
```

Track_artists Table:

```
MySQL 104.197.129.116:3306 ssl music SQL > SELECT COUNT(*) FROM track_artists;

+-----+
| COUNT(*) |
+-----+
|     12029 |
+-----+
1 row in set (0.0759 sec)
```

Advanced Queries & Top 15 Rows

- SELECT a.id, AVG(danceability), AVG(duration_ms), AVG(tempo) FROM tracks t JOIN artists a ON a.id = t.album_id GROUP BY t.album_id;

```
MySQL 104.197.129.116:3306 ssl music SQL > SELECT a.id, AVG(danceability), AVG(duration_ms), AVG(tempo)
-> FROM tracks t JOIN artists a ON a.id = t.album_id
-> GROUP BY t.album_id
-> LIMIT 15;
```

id	AVG(danceability)	AVG(duration_ms)	AVG(tempo)
0	0.540000	205906.5000	128.760000
1	0.681667	186159.6667	120.568333
2	0.660000	197933.0000	101.970000
3	0.540000	238413.0000	140.060000
4	0.640000	184539.5000	120.455000
5	0.750000	188490.0000	120.960000
6	0.665000	185640.0000	107.015000
7	0.360000	291719.5000	130.625000
8	0.405000	207873.0000	157.950000
9	0.810000	253886.0000	117.390000
10	0.580000	229426.0000	73.910000
11	0.670000	216399.5000	129.865000
12	0.390000	317835.3333	115.676667
13	0.680000	251866.0000	158.310000
14	0.530000	156640.0000	140.490000

15 rows in set (0.0291 sec)

- SELECT * FROM albums a JOIN (SELECT album_id, COUNT(id) as TotalTracks FROM tracks GROUP BY album_id) as total ON a.id = total.album_id WHERE TotalTracks > 1 AND (release_date LIKE '2009%' OR release_date LIKE '2005%') AND spotify_id LIKE '%a%';

```
MySQL 104.197.129.116:3306 ssl music SQL > SELECT * FROM albums a JOIN (SELECT album_id, COUNT(id) as TotalTracks FROM tracks GROUP BY album_id) as total
ON a.id = total.album_id WHERE TotalTracks > 1 AND (release_date LIKE '2009%' OR release_date LIKE '2005%') AND spotify_id LIKE '%a%' LIMIT 15;
```

id	name	spotify_id	album_id	TotalTracks	release_date	image_url
0	Conditions (Tour Edition)	6dfa3919e5acba68f0	0	2	2009-01-01	https://i.scdn.co/image/ab67616d0000b273f86ae8
7	X&Y	0V59MMTgoruvEqMv18KA0H	7	2	2005-06-07	https://i.scdn.co/image/ab67616d0000b2734e0362
141	Brand New Eyes	c225863f6ae2432651	141	5	2009-09-22	https://i.scdn.co/image/ab67616d0000b273b9abbe
160	21st Century Breakdown	dc516dd297039977bd	160	2	2009-05-15	https://i.scdn.co/image/ab67616d0000b273c2ced3
253	Music Of The Sun	9899b0d67cd5a724fa	253	2	2005-08-29	https://i.scdn.co/image/ab67616d0000b2734a8c86
515	Ladyhawke (Deluxe Edition)	9f3bd4d0d4519a8c50	515	3	2009-04-10	https://i.scdn.co/image/ab67616d0000b2730000e4
522	The Blueprint 3	7a4e869d4323ad0e3d	522	3	2009-09-08	https://i.scdn.co/image/ab67616d0000b2734b328c
614	Apocalypse	297d151d432c7b1aa3	614	4	2009-01-01	https://i.scdn.co/image/ab67616d0000b273331828
706	The Best of Dire Straits & Mark Knopfler - Private Investigations (Limited Edition)	2bfa1b2c5c7ed36274	706	3	2005-11-07	https://i.scdn.co/image/ab67616d0000b27311cd36
815	Gold	07f236dd71b56b1029	815	2	2005-01-01	https://i.scdn.co/image/ab67616d0000b273f24a70
914	Junior	6vQMBwthchxuSioACn2hcE	914	2	2009-03-18	https://i.scdn.co/image/ab67616d0000b273061dfb
920	La Roux	4c178ed97c2c1a0241	920	4	2009-01-01	https://i.scdn.co/image/ab67616d0000b2731ea480
947	This Is War	797874eeef69f2458	947	2	2009-01-01	https://i.scdn.co/image/ab67616d0000b27364219d
1061	The Emancipation of Mimi (Ultra Platinum Edition)	2be6dd466f96aa13a0	1061	4	2005-01-01	https://i.scdn.co/image/ab67616d0000b273923a02
1074	Monkey Business	29940be7edb73bfff87	1074	4	2005-01-01	https://i.scdn.co/image/ab67616d0000b27377234f

15 rows in set (0.0821 sec)

- SELECT a.id, t.name, t.duration_ms, t.danceability, t.loudness FROM albums a JOIN tracks t ON a.id = t.album_id WHERE t.duration_ms > 23000 and danceability > 0.70

and loudness > -3 UNION SELECT a2.id, t2.name, t2.duration_ms, t2.danceability, t2.loudness FROM albums a2 JOIN tracks t2 on a2.id = t2.album_id WHERE t2.duration_ms < 20000 and t2.danceability < 0.30 and t2.loudness < -8;

	id	name	duration_ms	danceability	loudness
▶	60	Too Much (feat. Usher)	165704	0.71	-2.76
	73	American Boy (feat. Kanye West)	284733	0.73	-2.99
	77	Truth Hurts	173306	0.71	-2.89
	256	Don't Wanna Let You Go - Radio Edit	217533	0.74	-2.90
	311	Great DJ	202813	0.79	-2.01
	335	Humpin' Around - Radio Edit	322413	0.74	-2.72
	379	Hey Ya!	235213	0.73	-2.26
	491	Brokenhearted	227146	0.77	-2.73
	309	Midnight Midnight	161213	0.74	-2.40
	1147	Get Up (Rattle) - Radio Edit	166932	0.80	-2.69
	1178	Put Your Hand Up - Radio Mix	208920	0.83	-2.11
	1179	2012 (It Ain't The End)	222200	0.72	-2.70
	866	He Don't Love You - Remastered	191760	0.74	-1.09
	1264	Angel	235133	0.74	-2.94
	1392	F.U.R.B. (F U Right Back)	201866	0.79	-2.92

- SELECT artists.name, AVG(tracks.energy), AVG(tracks.duration_ms) FROM (tracks JOIN track_artists ON tracks.id = track_artists.track_id) JOIN artists ON track_artists.artist_id = artists.id GROUP BY artists.name

name	AVG(tracks.energ...	AVG(tracks.duration_...
The Temper Trap	0.762857	226481.2857
Frankie Valli & The Four Seasons	0.551250	167491.3750
Foxes	0.790000	260421.6667
Captain & Tennille	0.650000	205986.2500
Rita Ora	0.785417	204359.3333
Coldplay	0.682500	267664.2500
Faith Hill	0.680000	207873.0000
The Police	0.577647	250306.9412
Chicago	0.652000	242709.0000
Urban Cookie Collective	0.875000	216399.5000
Guns N' Roses	0.775000	336476.5000
Rod Stewart	0.569286	278067.2143
We Five	0.530000	156640.0000
Christine Anu	0.795000	210579.5000
Vance Joy	0.746250	216650.0000
Guy Mitchell	0.560000	153901.0000
Wham!	0.742500	240958.1250
Tori Kelly	0.533333	193461.0000
Kungs	0.762500	191493.2500
The Monkees	0.655833	154714.2500

Part 2: Indexing

Command 1: SELECT a.id, AVG(danceability), AVG(duration_ms), AVG(tempo) FROM tracks t JOIN artists a ON a.id = t.album_id GROUP BY t.album_id;

Initial EXPLAIN ANALYZE w/ No Indexing:

```
MySQL 104.197.129.116:3306 ssl music SQL> EXPLAIN ANALYZE
-> SELECT a.id, AVG(danceability), AVG(duration_ms), AVG(tempo)
-> FROM tracks t JOIN artists a ON a.id = t.album_id
->
-> GROUP BY t.album_id;

+-----+
| EXPLAIN |
+-----+
|
+-----+
|
+-----+
| -> Table scan on <temporary> (actual time=24.1..24.6 rows=3812 loops=1)
|   -> Aggregate using temporary table (actual time=24.1..24.1 rows=3812 loops=1)
|     -> Nested loop inner join (cost=1977 rows=4616) (actual time=0.775..18.8 rows=5875 loops=1)
|       -> Covering index scan on a using PRIMARY (cost=361 rows=3554) (actual time=0.083..1.11 rows=3812 loops=1)
|       -> Index lookup on t using fk_tracks_albums (album_id=a.id) (cost=0.325 rows=1.3) (actual time=0.00345..0.00427 rows=1.54 loops=3812)
|
+-----+
|
+-----+
1 row in set (0.0824 sec)
```

- The cost is 1977 without Indexing.

EXPLAIN ANALYZE with index on tracks(danceability):

```
MySQL 104.197.129.116:3306 ssl music SQL> create index danceability on tracks(danceability);
Query OK, 0 rows affected (0.3130 sec)

Records: 0 Duplicates: 0 Warnings: 0
MySQL 104.197.129.116:3306 ssl music SQL> EXPLAIN ANALYZE SELECT a.id, AVG(danceability), AVG(duration_ms), AVG(tempo) FROM tracks t JOIN artists a ON
a.id = t.album_id GROUP BY t.album_id;

+-----+
| EXPLAIN |
+-----+
|
+-----+
|
+-----+
| -> Table scan on <temporary> (actual time=22.8..23.2 rows=3812 loops=1)
|   -> Aggregate using temporary table (actual time=22.8..22.8 rows=3812 loops=1)
|     -> Nested loop inner join (cost=1977 rows=4616) (actual time=0.0918..17.6 rows=5875 loops=1)
|       -> Covering index scan on a using PRIMARY (cost=361 rows=3554) (actual time=0.0547..1.08 rows=3812 loops=1)
|       -> Index lookup on t using fk_tracks_albums (album_id=a.id) (cost=0.325 rows=1.3) (actual time=0.00336..0.00412 rows=1.54 loops=3812)
|
+-----+
|
+-----+
1 row in set (0.0556 sec)
```

- The cost is 1977 with this index.

EXPLAIN ANALYZE with index on tracks(duration_ms):

Command 2: SELECT * FROM albums a JOIN (SELECT album_id, COUNT(id) as TotalTracks FROM tracks GROUP BY album_id) as total ON a.id = total.album_id WHERE TotalTracks > 1 AND (release_date LIKE '2009%' OR release_date LIKE '2005%') AND spotify_id LIKE '%a%';

Initial EXPLAIN ANALYZE w/ No Indexing:

```
EXPLAIN ANALYZE
SELECT * FROM albums a
JOIN (SELECT album_id, COUNT(id) as TotalTracks
FROM tracks GROUP BY album_id) as total
ON a.id = total.album_id
WHERE TotalTracks > 1 AND (release_date LIKE '2009%' OR release_date LIKE '2005%') AND spotify_id LIKE '%a%';
```

r | Navigate: ⏮ ⏪ ⏩ ⏭ |

```
-> Nested loop inner join (cost=126299 rows=1.25e+6) (actual time=6.01..12.7 rows=59 loops=1)
-> Filter: (((a.release_date like '2009%') or (a.release_date like '2005%')) and (a.spotify_id like '%a%')) (cost=743 rows=168) (actual time=0.127..6.54 rows=224 loops=1)
-> Table scan on a (cost=743 rows=7185) (actual time=0.117..3.7 rows=7470 loops=1)
-> Index lookup on total using <auto_key0> (album_id=a.id) (cost=2697..2700 rows=10) (actual time=0.0272..0.0273 rows=0.263 loops=224)
```

- Cost: 126,999 w/o indexing

```
create index rel
on albums(release_date);
EXPLAIN ANALYZE
SELECT * FROM albums a
JOIN (SELECT album_id, COUNT(id) as TotalTracks
FROM tracks GROUP BY album_id) as total
ON a.id = total.album_id
WHERE TotalTracks > 1 AND (release_date LIKE '2009%' OR release_date LIKE '2005%') AND spotify_id LIKE '%a%';
```

r | Navigate: ⏮ ⏪ ⏩ ⏭ |

```
-> Nested loop inner join (cost=126299 rows=1.25e+6) (actual time=5.98..12.5 rows=59 loops=1)
-> Filter: (((a.release_date like '2009%') or (a.release_date like '2005%')) and (a.spotify_id like '%a%')) (cost=743 rows=168) (actual time=0.0842..6.33 rows=224 loops=1)
-> Table scan on a (cost=743 rows=7185) (actual time=0.0757..3.51 rows=7470 loops=1)
-> Index lookup on total using <auto_key0> (album_id=a.id) (cost=2697..2700 rows=10) (actual time=0.0272..0.0272 rows=0.263 loops=224)
```

- Cost: 126,999 w/ indexing on albums.release_date

```
create index rel
on albums(spotify_id);
EXPLAIN ANALYZE
SELECT * FROM albums a
JOIN (SELECT album_id, COUNT(id) as TotalTracks
FROM tracks GROUP BY album_id) as total
ON a.id = total.album_id
WHERE TotalTracks > 1 AND (release_date LIKE '2009%' OR release_date LIKE '2005%') AND spotify_id LIKE '%a%';
```

r | Navigate: ⏮ ⏪ ⏩ ⏭ |

```
-> Nested loop inner join (cost=126299 rows=1.25e+6) (actual time=6.04..12.8 rows=59 loops=1)
-> Filter: (((a.release_date like '2009%') or (a.release_date like '2005%')) and (a.spotify_id like '%a%')) (cost=743 rows=168) (actual time=0.0955..6.54 rows=224 loops=1)
-> Table scan on a (cost=743 rows=7185) (actual time=0.0874..3.67 rows=7470 loops=1)
-> Index lookup on total using <auto_key0> (album_id=a.id) (cost=2697..2700 rows=10) (actual time=0.0275..0.0276 rows=0.263 loops=224)
```

- Cost: 126,999 w/ indexing on albums.spotify_id

```

create index spot
on albums(release_date, spotify_id);
EXPLAIN ANALYZE
SELECT * FROM albums a
JOIN (SELECT album_id, COUNT(id) as TotalTracks
FROM tracks GROUP BY album_id) as total
ON a.id = total.album_id
WHERE TotalTracks > 1 AND (release_date LIKE '2009%' OR release_date LIKE '2005%') AND spotify_id LIKE '%a%';

```

· | Navigate: ⏮ ⏪ ⏩ ⏭ |

```

-> Nested loop inner join (cost=126299 rows=1.25e+6) (actual time=5.33..12.3 rows=59 loops=1)
  -> Filter: (((a.release_date like '2009%') or (a.release_date like '2005%')) and (a.spotify_id like '%a%')) (cost=743 rows=168) (actual time=0.093..6.79 rows=224 loops=1)
  -> Table scan on a (cost=743 rows=7185) (actual time=0.0846..3.74 rows=7470 loops=1)
  -> Index lookup on total using <auto_key0> (album_id=a.id) (cost=2697..2700 rows=10) (actual time=0.0243..0.0244 rows=0.263 loops=224)

```

- Cost: 126,999 w/ indexing on albums.spotify_id and albums.release_date

In this query, we were unable to reduce the cost in every implementation of indexing. This is likely due to the fact that the 'WHERE' condition includes the '%' symbol which represents a string "wildcard". As a result, the query must examine every single row in the database regardless of any indexing. For instance, in the condition "spotify_id LIKE '%a'", the database must scan all values to see which rows contain the character 'a'. Therefore, indexing did not improve the cost of our query.

Command 3: SELECT a.id, t.name, t.duration_ms, t.danceability, t.loudness FROM albums a JOIN tracks t ON a.id = t.album_id WHERE t.duration_ms > 23000 and danceability > 0.70 and loudness > -3 UNION SELECT a2.id, t2.name, t2.duration_ms, t2.danceability, t2.loudness FROM albums a2 JOIN tracks t2 on a2.id = t2.album_id WHERE t2.duration_ms < 20000 and t2.danceability < 0.30 and t2.loudness < -8;

EXPLAIN ANALYZE w/o index

```

EXPLAIN ANALYZE
SELECT a.id, t.name, t.duration_ms, t.danceability, t.loudness FROM albums a
JOIN tracks t ON a.id = t.album_id
WHERE t.duration_ms > 23000 and danceability > 0.70 and loudness > -3
UNION
SELECT a2.id, t2.name, t2.duration_ms, t2.danceability, t2.loudness FROM albums a2
JOIN tracks t2 on a2.id = t2.album_id
WHERE t2.duration_ms < 20000 and t2.danceability < 0.30 and t2.loudness < -8;

```

r | Navigate: ⏮ ⏪ ⏩ ⏭ |

```

-> Table scan on <union temporary> (cost=2322..2333 rows=701) (actual time=13.1..13.1 rows=96 loops=1)
  -> Union materialize with deduplication (cost=2322..2322 rows=701) (actual time=13.1..13.1 rows=96 loops=1)
  -> Nested loop inner join (cost=1126 rows=351) (actual time=0.135..7.43 rows=96 loops=1)
    -> Filter: ((t.duration_ms > 23000) and (t.danceability > 0.70) and (t.loudness > -3.00) and (t.album_id is not null)) (cost=1003 rows=351) (actual time=0.114..6.98

```

- Cost: 2322 w/o indexing


```

create index dur
on tracks(duration_ms);
EXPLAIN ANALYZE
SELECT a.id, t.name, t.duration_ms, t.danceability, t.loudness FROM albums a
JOIN tracks t ON a.id = t.album_id
WHERE t.duration_ms > 23000 and danceability > 0.70 and loudness > -3
UNION
SELECT a2.id, t2.name, t2.duration_ms, t2.danceability, t2.loudness FROM albums a2
JOIN tracks t2 on a2.id = t2.album_id
WHERE t2.duration_ms < 20000 and t2.danceability < 0.30 and t2.loudness < -8;

```

Navigation: ⏮ ⏪ 1 / 1 ⏩ ⏭

```

-> Table scan on <union temporary> (cost=1477..1493 rows=1052) (actual time=8.02..8.04 rows=96 loops=1)
-> Union materialize with deduplication (cost=1477..1477 rows=1052) (actual time=8.02..8.02 rows=96 loops=1)
-> Nested loop inner join (cost=1371 rows=1052) (actual time=0.172..7.85 rows=96 loops=1)
-> Filter: ((t.duration_ms > 23000) and (t.danceability > 0.70) and (t.loudness > -3.00) and (t.album_id is not null)) (cost=1003 rows=1052) (actual time=0.157..7.43

```

- Cost: 1477 with indexing on tracks.duration_ms

```

EXPLAIN ANALYZE
SELECT a.id, t.name, t.duration_ms, t.danceability, t.loudness FROM albums a
JOIN tracks t ON a.id = t.album_id
WHERE t.duration_ms > 23000 and danceability > 0.70 and loudness > -3
UNION
SELECT a2.id, t2.name, t2.duration_ms, t2.danceability, t2.loudness FROM albums a2
JOIN tracks t2 on a2.id = t2.album_id
WHERE t2.duration_ms < 20000 and t2.danceability < 0.30 and t2.loudness < -8;

```

Navigation: ⏮ ⏪ ⏩ ⏭

```

-> Table scan on <union temporary> (cost=1584..1598 rows=880) (actual time=7.49..7.51 rows=96 loops=1)
-> Union materialize with deduplication (cost=1584..1584 rows=880) (actual time=7.49..7.49 rows=96 loops=1)
-> Nested loop inner join (cost=1496 rows=880) (actual time=0.0567..7.32 rows=96 loops=1)
-> Filter: ((t.duration_ms > 23000) and (t.loudness > -3.00) and (t.album_id is not null)) (cost=1188 rows=880) (actual time=0.0442..6.87 rows=96 loops=1)

```

- Cost: 1584 with indexing on tracks.duration_ms and tracks.danceability

```

EXPLAIN ANALYZE
SELECT a.id, t.name, t.duration_ms, t.danceability, t.loudness FROM albums a
JOIN tracks t ON a.id = t.album_id
WHERE t.duration_ms > 23000 and danceability > 0.70 and loudness > -3
UNION
SELECT a2.id, t2.name, t2.duration_ms, t2.danceability, t2.loudness FROM albums a2
JOIN tracks t2 on a2.id = t2.album_id
WHERE t2.duration_ms < 20000 and t2.danceability < 0.30 and t2.loudness < -8;

```

Navigation: ⏮ ⏪ ⏩ ⏭

```

-> Table scan on <union temporary> (cost=219..223 rows=121) (actual time=2.12..2.14 rows=96 loops=1)
-> Union materialize with deduplication (cost=219..219 rows=121) (actual time=2.12..2.12 rows=96 loops=1)
-> Nested loop inner join (cost=207 rows=121) (actual time=0.0653..1.87 rows=96 loops=1)
-> Filter: ((t.duration_ms > 23000) and (t.danceability > 0.70) and (t.album_id is not null)) (cost=164 rows=121) (actual time=0.0518..1.61 rows=96 loops=1)

```

- Cost: 219 with indexing on tracks.duration_ms and tracks.loudness

Indexing Analysis:

For the final indexing design for this query, we chose to go with indexing on two columns, tracks.duration_ms and tracks.loudness. When running the command with EXPLAIN ANALYZE, we achieved a cost of 219.

In our first implementation of indexing, we noticed that the cost dramatically decreased from the original query (w/o indexing) and this made sense since it helped the query have better lookup times when filtering out rows. So, in an attempt to further reduce the query search time, for our second implementation of indexing, we indexed two columns: tracks.duration_ms and tracks.danceability. However, this implementation only increased the cost compared to a single indexing of tracks.duration_ms. One reason this might be is because this format of indexing reduced the index's selectivity, thus causing the query to examine more rows than the prior.

Finally, for our third implementation of indexing, we chose to index tracks.duration_ms and tracks.loudness. This turned out to be the most effective for filtering/sorting through the multiple columns. This is most likely because this method reduced the most rows that the query needed to examine before returning the results.

Command 4:

Initial EXPLAIN ANALYZE w/ No Indexing:

```
1 • EXPLAIN ANALYZE SELECT artists.name, AVG(tracks.energy), AVG(tracks.duration_ms)
2 FROM (tracks JOIN track_artists ON tracks.id = track_artists.track_id)
3 JOIN artists ON track_artists.artist_id = artists.id
4 GROUP BY artists.name
```

100% 1:3

Form Editor Navigate: 1 / 1

EXPLAIN:

- > Table scan on <temporary> (actual time=58.5..59.4 rows=3799 loops=1)
- > Aggregate using temporary table (actual time=58.5..58.5 rows=3799 loops=1)
- > Nested loop inner join (cost=6244 rows=11091) (actual time=0.152..42.9 rows=12029 loops=1)
- > Nested loop inner join (cost=2362 rows=11091) (actual time=0.14..13.7 rows=12029 loops=1)

- **Cost: 6244 without indexing**

EXPLAIN ANALYZE with index on artists(name)

```
1  EXPLAIN ANALYZE SELECT artists.name, AVG(tracks.energy), AVG(tracks.duration_ms)
2  FROM (tracks JOIN track_artists ON tracks.id = track_artists.track_id)
3  JOIN artists ON track_artists.artist_id = artists.id
4  GROUP BY artists.name
```

100% 1:4

Form Editor Navigate: 1/1

EXPLAIN:

- > Group aggregate: avg(tracks.energy), avg(tracks.duration_ms) (cost=7353 rows=3554) (actual time=0.158..49 rows=3799 loops=1)
- > Nested loop inner join (cost=6244 rows=11091) (actual time=0.132..42.7 rows=12029 loops=1)
- > Nested loop inner join (cost=2362 rows=11091) (actual time=0.108..15.1 rows=12029 loops=1)
- > Covering index scan on artists using foo (cost=361 rows=3554) (actual time=0.0681..1.25 rows=3862 loops=1)

- **Cost: 6244**

EXPLAIN ANALYZE with index on tracks(energy)

```
3  EXPLAIN ANALYZE SELECT artists.name, AVG(tracks.energy), AVG(tracks.duration_ms)
4  FROM (tracks JOIN track_artists ON tracks.id = track_artists.track_id)
5  JOIN artists ON track_artists.artist_id = artists.id
6  GROUP BY artists.name;
```

100% 17:3

Form Editor Navigate: 1/1

EXPLAIN:

- > Table scan on <temporary> (actual time=115..116 rows=3799 loops=1)
- > Aggregate using temporary table (actual time=115..115 rows=3799 loops=1)
- > Nested loop inner join (cost=6257 rows=11091) (actual time=3.29..95.2 rows=12029 loops=1)
- > Nested loop inner join (cost=2375 rows=11091) (actual time=2.59..67.3 rows=12029 loops=1)

- **Cost: 6257**

EXPLAIN ANALYZE with index on tracks(duration_ms)

```
3  EXPLAIN ANALYZE SELECT artists.name, AVG(tracks.energy), AVG(tracks.duration_ms)
4  FROM (tracks JOIN track_artists ON tracks.id = track_artists.track_id)
5  JOIN artists ON track_artists.artist_id = artists.id
6  GROUP BY artists.name
```

100% 1:6

Form Editor Navigate: 1/1

EXPLAIN:

- > Table scan on <temporary> (actual time=53.2..54.1 rows=3799 loops=1)
- > Aggregate using temporary table (actual time=53.2..53.2 rows=3799 loops=1)
- > Nested loop inner join (cost=6244 rows=11091) (actual time=0.104..39 rows=12029 loops=1)
- > Nested loop inner join (cost=2362 rows=11091) (actual time=0.0899..13 rows=12029 loops=1)

- **Cost: 6244**

Indexing Analysis: We chose to use indexing by the artist name as our final design. We tried indexing each attribute (and combinations of said attributes), and there was no improvement in

cost. We believe this is because the query itself is complex enough, so using an index does not improve the performance much, if at all. In addition, the index might not narrow down the result set significantly. In addition, the index may require a lot of maintenance.
